

Appliance Energy Consumption Forecasting Using Classical Time Series, Machine Learning and Foundation Models

GitHub Link: <https://github.com/Lojana96/Appliance-energy-forecasting>

Student Name: Lojana Jegatheeswaran

Student ID: 24106862

1. Introduction

The rapid growth of smart home technologies and intelligent energy management systems has increased the demand for accurate short-term electricity consumption forecasting. Reliable energy forecasts enable households and utility providers to optimise energy usage, reduce operational costs, improve demand-side management and support the efficient integration of renewable energy resources into modern power systems. As energy consumption exhibits complex temporal patterns influenced by behavioural and environmental factors, developing robust forecasting models has become an important research area within data analytics and artificial intelligence (Hyndman and Athanasopoulos, 2021).

This study investigates the effectiveness of different forecasting approaches for predicting household appliance energy consumption over a 24-hour forecasting horizon. The analysis follows a systematic forecasting framework consisting of data preparation, exploratory time-series analysis, benchmark forecasting, statistical modelling, feature-based machine learning and foundation models. The forecasting approaches are evaluated using standard performance metrics, including Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), Mean Absolute Scaled Error (MASE) and forecast bias. The findings are used to compare forecasting performance and identify the most appropriate approach by considering predictive accuracy, computational efficiency and practical applicability for residential energy forecasting.

2. Dataset Description and Preprocessing

2.1 Dataset Description

This study uses the **Appliance Energy Prediction** dataset, which is publicly available through the **UCI Machine Learning Repository** and was originally developed by Candanedo, Feldheim and

Deramaix (2017). The dataset was collected from a low-energy residential house in Belgium over approximately four and a half months, from January to May 2016, to investigate the relationship between household appliance energy consumption and indoor environmental conditions. Measurements were recorded at **10-minute intervals**, producing **19,735 observations**. The dataset includes appliance energy consumption together with indoor temperature and relative humidity measurements collected from multiple rooms of the house, as well as outdoor weather variables such as temperature, humidity, atmospheric pressure, wind speed, visibility and dew point. These variables provide valuable contextual information for modelling household energy consumption because appliance usage is influenced by both indoor environmental conditions and external weather factors (Candanedo *et al.*, 2017).

2.2 Data Preprocessing

Before model development, several preprocessing steps were performed to prepare the dataset for time-series forecasting. The timestamp variable was converted into a datetime format and used as the time-series index to preserve the chronological order of observations. The original 10-minute observations were resampled into **hourly observations**, reducing the dataset from **19,735** to **3,290 observations**. This resampling process reduced short-term fluctuations while preserving the underlying temporal patterns and significantly improved the computational efficiency of subsequent forecasting models. Data quality checks confirmed that the processed dataset contained no missing values after resampling. The target variable used throughout this study is **Appliances**, representing household appliance energy consumption, while the remaining environmental measurements were treated as potential explanatory variables for statistical and machine learning forecasting models.

3. Exploratory Data Analysis

3.1 Time Series Characteristics




Figure 1: Hourly Appliance Energy Consumption Time Series (January–May 2016)

Figure 1 presents the hourly appliance energy consumption after aggregating the original 10-minute observations into hourly intervals. The time series exhibits substantial variability, with consumption ranging from approximately **28 Wh** to **608 Wh**, indicating frequent fluctuations in household electricity demand. Although no pronounced long-term upward or downward trend is evident over the five-month observation period, recurring peaks and troughs suggest the presence of regular temporal patterns associated with household activities.

The repeated high-demand periods indicate that appliance usage follows a structured daily cycle rather than random behaviour. These characteristics suggest that the data contain temporal dependencies and seasonal effects, making them suitable for time-series forecasting methods such as SARIMAX and machine learning models capable of capturing sequential patterns.




Figure 2: Decomposition of the appliance energy consumption time series, illustrating the underlying trend, recurring daily seasonal pattern,

and unexplained residual variation used to assess temporal structure prior to forecasting.

Figure 2 illustrates the seasonal decomposition of the hourly appliance energy consumption into its **observed**, **trend**, **seasonal**, and **residual** components. The trend component shows moderate fluctuations in average energy usage throughout the observation period without a persistent long-term increase or decrease. In contrast, the seasonal component exhibits a strong and highly regular repeating pattern every 24 hours, indicating that appliance consumption follows a consistent daily cycle driven by household activities.

The residual component is centred around zero but contains several large positive and negative deviations, suggesting that occasional irregular events and unexpected appliance usage are not fully explained by the trend and seasonal effects. Overall, the decomposition confirms that the series contains strong daily seasonality together with random variation, supporting the selection of seasonal forecasting models such as **SARIMAX** with a 24-hour seasonal period.




Figure 3: Autocorrelation (ACF) and partial autocorrelation (PACF) analysis of hourly appliance energy consumption, showing significant temporal dependence and a repeating 24-hour seasonal cycle used to guide SARIMAX parameter selection.

Figure 3 presents the autocorrelation function (ACF) and partial autocorrelation function (PACF) of the hourly appliance energy consumption series. The ACF displays significant positive autocorrelation at short lags and pronounced peaks at multiples of **24 hours**, confirming the presence of strong daily seasonality. The gradual decay in autocorrelation further indicates that current energy consumption is influenced by previous observations over multiple time steps.

The PACF exhibits significant spikes at the first few lags followed by a rapid decline, suggesting that only a small number of autoregressive terms are required to model the short-term dependence. Together, the ACF and PACF plots provide statistical evidence for incorporating both non-seasonal and seasonal autoregressive components within the SARIMAX model, while confirming that the series possesses strong temporal dependence suitable for time-series forecasting.

ADF Test Results		KPSS Test Results	
series	Original hourly Appliances	series	Original hourly Appliances
ADF Statistic	-8.948888	KPSS Statistic	0.038399
p-value	0.0	p-value	0.1
Used Lags	29	Used Lags	25
Observations	3260	Critical Value 10%	0.347
Critical Value 1%	-3.432358	Critical Value 5%	0.463
Critical Value 5%	-2.862427	Critical Value 2.5%	0.574
Critical Value 10%	-2.567242	Critical Value 1%	0.739
dtype: object		dtype: object	

Figure 4: Results of the Augmented Dickey–Fuller (ADF) and Kwiatkowski–Phillips–Schmidt–Shin (KPSS) stationarity tests used to assess the statistical properties of the hourly appliance energy consumption series before model development.

Figure 4 presents the results of the complementary **ADF** and **KPSS** stationarity tests. The **ADF test** produced a test statistic of **−8.95** with a **p-value < 0.001**, leading to rejection of the null hypothesis of a unit root. Conversely, the **KPSS test** returned a statistic of **0.038** with a **p-value of 0.10**, indicating insufficient evidence to reject the null hypothesis of stationarity. The agreement between these two tests provides strong statistical evidence that the hourly appliance energy consumption series is stationary.

These findings indicate that the underlying statistical properties of the series remain relatively stable over time, allowing autoregressive forecasting models to be fitted without extensive non-seasonal differencing. Consequently, the SARIMAX modelling process focused primarily on capturing the strong daily seasonal behaviour identified during the exploratory analysis.

4. Forecasting Methodology

The forecasting task was defined as predicting appliance energy consumption over a 24-hour horizon using hourly observations. A chronological split was applied to preserve the temporal ordering of the series, with the final 14 days (336

observations) reserved as the test period. Random train-test splitting was avoided because it would allow future information to influence model development and would therefore be inappropriate for time-series forecasting. All models were evaluated using the same accuracy measures: MAE, RMSE, MASE and forecast bias, enabling consistent comparison across simple and advanced forecasting approaches (Hyndman and Athanasopoulos, 2021).

4.1 Benchmark Models

Five benchmark forecasting methods were implemented: Mean, Naïve, Daily Seasonal Naïve, Weekly Seasonal Naïve and Drift. Benchmark models provide a necessary reference for assessing whether more complex forecasting techniques generate genuine improvements rather than merely increasing model complexity. The daily seasonal naïve model used a seasonal period of 24 hours, while the weekly seasonal naïve model used 168 hours, corresponding to seven days of hourly observations. These benchmarks were selected because they represent different assumptions about the underlying temporal structure, ranging from constant average demand to persistence, daily repetition and weekly repetition (Hyndman and Athanasopoulos, 2021).




Figure 5: Comparison of benchmark forecasting models using Root Mean Squared Error (RMSE), demonstrating the relative predictive accuracy of the baseline forecasting approaches over the test period.

Figure 5 compares the predictive performance of the five benchmark forecasting models using the Root Mean Squared Error (RMSE). The **Mean model** achieved the lowest forecasting error (**RMSE = 74.91**), indicating the strongest overall baseline performance. The **Weekly Seasonal Naïve model** produced a comparable result

(RMSE = **79.29**), suggesting that weekly repetition captures some of the underlying consumption behaviour. In contrast, the **Daily Seasonal Naïve model** resulted in a considerably higher error (RMSE = **129.23**), while the **Naïve (258.82)** and **Drift (274.61)** models exhibited the weakest predictive performance.

These findings demonstrate that simple persistence-based forecasting methods are insufficient for accurately modelling the complex variability of household appliance energy consumption. Consequently, more sophisticated forecasting approaches, including **SARIMAX**, feature-based machine learning models, and foundation models, are expected to achieve improved forecasting accuracy by modelling both temporal dependence and seasonal behaviour.

	Model	MAE	RMSE	MASE	Bias	RMSE_Improvement_vs_Naive_%
0	Mean	50.319	74.906	0.942	-3.109	71.059
1	Weekly Seasonal Naive	42.634	79.290	0.798	-10.818	69.365
2	Daily Seasonal Naive	86.959	129.232	1.628	64.013	50.069
3	Naive	250.640	258.820	4.692	247.763	0.000
4	Drift	266.373	274.611	4.986	264.501	-6.101

Figure 6:Performance comparison of benchmark forecasting models on the 14-day test dataset.

Figure 6 summarises the performance of the benchmark forecasting models. The Mean model achieved the lowest RMSE (74.91), while the Weekly Seasonal Naïve model recorded the lowest MAE (42.63) and MASE (0.798), indicating competitive forecasting performance. In contrast, the Naïve and Drift models produced substantially larger forecasting errors, confirming that simple persistence-based approaches are insufficient for accurately modelling household electricity demand. These benchmark results provide a performance baseline against which the SARIMAX and machine learning models are evaluated.

4.2 SARIMAX Model

The Seasonal AutoRegressive Integrated Moving Average with eXogenous variables (SARIMAX) model was implemented as the primary statistical forecasting model to capture both temporal dependence and daily seasonality in household appliance energy consumption. A two-stage parameter optimisation procedure was adopted. First, a non-seasonal grid search identified the optimal autoregressive, differencing and moving

average parameters using the minimum Akaike Information Criterion (AIC). The selected non-seasonal order was (1, 0, 6). Subsequently, a constrained seasonal grid search was performed using a seasonal period of 24 hours, resulting in an optimal seasonal order of (1, 1, 1, 24). The final model achieved an AIC of 32323.314 and BIC of 32442.757, indicating an appropriate balance between model complexity and goodness of fit.




Figure 7: Comparison of observed and predicted hourly appliance energy consumption using the final SARIMAX model with 95% prediction intervals over the 24-hour forecast horizon.

Figure 7 compares the observed appliance energy consumption with the forecasts generated by the final SARIMAX model over the 24-hour prediction period. The model successfully captures the overall daily consumption profile and follows the observed trend throughout most of the forecast horizon. Although several sharp consumption peaks are slightly underestimated, the majority of the observed values remain within the 95% prediction intervals, indicating that the model provides reliable uncertainty estimates. These results demonstrate that the selected SARIMAX model effectively captures the strong daily seasonal behaviour identified during the exploratory data analysis and substantially improves forecasting performance compared with the benchmark models.

	Model	MAE	RMSE	MASE	Bias
0	SARIMAX	36.71	65.135	0.687	-9.081

Figure 8:SARIMAX evaluation table.

The rolling-origin evaluation demonstrates that the SARIMAX model substantially improves forecasting accuracy compared with all benchmark models. The model achieved an RMSE of **65.14**,

representing approximately a **13% improvement** over the best benchmark model (RMSE = 74.91). Furthermore, the MASE value below one confirms that the SARIMAX model consistently outperformed the naïve forecasting baseline while maintaining only a small negative forecasting bias.

4.3 Feature-Based Machine Learning Model

To investigate whether engineered predictor variables could further improve forecasting accuracy, a HistGradientBoosting regression model was developed using calendar features together with lagged and rolling statistical features derived from the historical appliance energy series. Gradient boosting was selected because of its ability to model complex non-linear relationships while remaining computationally efficient for medium-sized datasets. The model was trained using the same chronological train-test split as the statistical models to ensure a fair comparison.




Figure 9:Comparison of observed and predicted hourly appliance energy consumption using the HistGradientBoosting feature-based forecasting model over the 24-hour prediction horizon

Figure 9 illustrates the forecasting performance of the HistGradientBoosting model over the 24-hour evaluation period. The predicted values closely follow the observed appliance energy consumption and successfully capture the overall daily demand pattern. Compared with the benchmark models, the feature-based model provides smoother predictions while accurately modelling both gradual changes and short-term fluctuations in electricity consumption. Although a small number of extreme demand peaks are underestimated, the model demonstrates strong predictive capability, highlighting the effectiveness of engineered temporal features for short-term residential energy forecasting.

	Model	MAE	RMSE	MASE	Bias
0	HistGradientBoosting	30.956	52.255	0.6	3.154

Figure 10:Forecasting Performance of the HistGradientBoosting Feature-Based Model

Figure 10 summarises the forecasting performance of the HistGradientBoosting model over the rolling evaluation period. The model achieved an RMSE of **52.26**, improving upon both the benchmark models and the SARIMAX model. The MASE value of **0.600** indicates that the model consistently outperformed the naïve forecasting baseline, while the small positive bias (**3.15**) suggests only a slight tendency to overestimate appliance energy consumption. These results demonstrate that feature engineering combined with gradient boosting provides an effective approach for modelling the non-linear characteristics of household electricity demand.

4.4 Foundation Model

Recent advances in foundation models have demonstrated promising zero-shot forecasting capabilities for time-series prediction. In this study, the Chronos-Bolt-Tiny foundation model was evaluated without task-specific fine-tuning to assess whether a pre-trained forecasting model could accurately predict household appliance energy consumption. Unlike conventional statistical and machine learning approaches, Chronos relies on knowledge learned from large collections of time-series datasets rather than handcrafted feature engineering.




Figure 11: Comparison of observed and predicted hourly appliance energy consumption using the Chronos-Bolt-Tiny foundation forecasting model over the 24-hour prediction horizon

Figure 11 illustrates the forecasting performance of the Chronos-Bolt-Tiny foundation model over the 24-hour evaluation period. The model accurately

captures the overall daily consumption trend and produces smooth forecasts across the prediction horizon. Although several high-demand periods are overestimated, the model achieves the lowest forecasting error among all evaluated approaches, demonstrating the effectiveness of pretrained transformer-based foundation models for short-term household energy forecasting without extensive model-specific training.

	Model	MAE	RMSE	MASE	Bias
0	Chronos-Bolt-Tiny	29.828609	42.225821	0.55834	16.751496

Figure 12:Forecasting Performance of the Chronos-Bolt-Tiny Foundation Model

Figure 12 presents the forecasting performance of the Chronos-Bolt-Tiny foundation model. The model achieved the lowest RMSE (**42.23**) and MAE (**29.83**) among all evaluated forecasting approaches, indicating the highest overall prediction accuracy. The MASE value of **0.558** confirms that the model substantially outperformed the naïve forecasting baseline. Although the positive bias (**16.75**) indicates a slight tendency to overestimate energy consumption during some high-demand periods, the foundation model consistently delivered the best forecasting performance in this study.

4.5 Final Model Comparison




Figure 13:Comparison of forecasting accuracy (RMSE) for the evaluated forecasting models over the common 24-hour test period.

Figure 13 compares the forecasting accuracy of the four evaluated models using the Root Mean Squared Error (RMSE). Lower RMSE values indicate better forecasting performance. The SARIMAX model achieved the lowest RMSE (24.16), followed by the HistGradientBoosting model (29.87), while the Chronos-Bolt-Tiny

foundation model obtained an RMSE of 42.23. The Weekly Seasonal Naïve benchmark produced the highest RMSE (48.81), confirming that all proposed forecasting models outperformed the strongest benchmark model.

	Model	MAE	RMSE	MASE	Bias
0	SARIMAX	18.841	24.158	0.353	-7.402
1	HistGradientBoosting	19.360	29.865	0.362	2.077
2	Chronos-Bolt-Tiny	29.829	42.226	0.558	16.751
3	Weekly Seasonal Naïve	30.000	48.808	0.562	-6.528

Figure 14:Comparative forecasting performance of the final forecasting models over the common 24-hour evaluation period.

Figure 14 compares the forecasting performance of the final statistical, machine learning, foundation, and benchmark forecasting models using a common 24-hour evaluation period. The SARIMAX model achieved the highest overall forecasting accuracy, recording the lowest RMSE (24.16), MAE (18.84), and MASE (0.353). The HistGradientBoosting model ranked second with an RMSE of 29.87, followed by the Chronos-Bolt-Tiny foundation model (42.23). Although the Weekly Seasonal Naïve benchmark performed reasonably well, all three proposed forecasting approaches produced lower forecasting errors, demonstrating the benefits of advanced forecasting techniques over traditional benchmark methods.

5. Discussion and Analysis

5.1 Which benchmark model is strongest?

The **Weekly Seasonal Naïve** model achieved the lowest forecasting error among the benchmark models (RMSE = 79.29), outperforming the Mean, Naïve, Daily Seasonal Naïve and Drift models. This indicates that household appliance energy consumption exhibits a strong **weekly seasonal pattern**, where electricity usage during a given hour is more similar to the corresponding hour in the previous week than to the previous day. The poor performance of the Naïve and Drift models further suggests that the series is dominated by recurring seasonal behaviour rather than a simple trend.

5.2 Does SARIMAX improve on the strongest benchmark?

Yes. The SARIMAX model substantially improved upon the strongest benchmark. During rolling evaluation, SARIMAX reduced RMSE from **79.29** (Weekly Seasonal Naïve) to **65.14**, while in the common 24-hour comparison it achieved the lowest RMSE (**24.16**). This demonstrates that SARIMAX successfully captures both short-term autocorrelation and the dominant daily seasonal cycle. The seasonal order (1,1,1,24) effectively models the 24-hour periodicity observed during exploratory analysis. Although no additional exogenous variables were incorporated, the model still produced highly accurate forecasts by exploiting the temporal structure of the series.

5.3 Does the feature-based model improve after adding engineered features?

Yes. The HistGradientBoosting model achieved an RMSE of **52.26**, improving upon both the benchmark models and SARIMAX in the rolling evaluation. The inclusion of lag features, rolling-window statistics, hour-of-day, day-of-week and calendar features enabled the model to learn complex non-linear relationships that are difficult to capture using traditional statistical models. Among these, lag variables and rolling-window statistics appeared to contribute most to forecasting accuracy because they directly represent recent consumption behaviour.

5.4 Does the foundation model outperform the others?

The Chronos-Bolt-Tiny foundation model achieved an RMSE of 42.23 during the 24-hour evaluation, outperforming the feature-based model and the benchmark model but not the SARIMAX model in the common comparison. Although the foundation model required minimal task-specific training and demonstrated competitive performance, the improvement was not sufficiently large to justify its additional computational complexity for this dataset. Consequently, pretrained foundation models provide a promising alternative, but simpler statistical or machine learning models may remain preferable for practical deployment on relatively small datasets.

5.5 Which variables are genuinely known at the forecast origin?

Variables such as historical appliance energy consumption, lag features, rolling statistics, hour of day, day of week and calendar information are all available at the forecast origin and therefore represent genuine forecasting inputs. However, future indoor temperature, humidity or weather observations would not normally be available when producing operational forecasts. If such future observations are used directly from the test dataset, the prediction becomes a **conditional forecast** rather than a true forecast because it assumes future explanatory variables are already known.

5.6 Which model would you recommend?

Considering forecasting accuracy, interpretability, computational efficiency and ease of deployment, the **SARIMAX model** is recommended for practical smart-home energy forecasting. It achieved the best overall performance during the common evaluation period (RMSE = **24.16**) while remaining computationally efficient and highly interpretable. Although the HistGradientBoosting and Chronos-Bolt-Tiny models also produced accurate forecasts, they require greater computational resources and more complex implementation. Therefore, SARIMAX provides the best balance between predictive accuracy, reliability and deployment practicality for short-term residential energy forecasting.

References

- Candanedo, L.M., Feldheim, V. and Deramaix, D. (2017) 'Data driven prediction models of energy use of appliances in a low-energy house', *Energy and Buildings*, 140, pp. 81–97.
- Dua, D. and Graff, C. (2019) *UCI Machine Learning Repository*. University of California, Irvine. Available at: <https://archive.ics.uci.edu/ml> (Accessed: 13 August 2026).
- Hyndman, R.J. and Athanasopoulos, G. (2021) *Forecasting: Principles and Practice*. 3rd edn. Melbourne: OTexts.